

EE4210: Computer Communication Networks II

Socket Programming



Outline

- Principles of socket programming
- Example: TCP based client-server

Basics

- Transport layer: TCP/UDP
- C programming basics
- Data structures basics
- Windows: winsock library

Telephone Analogy

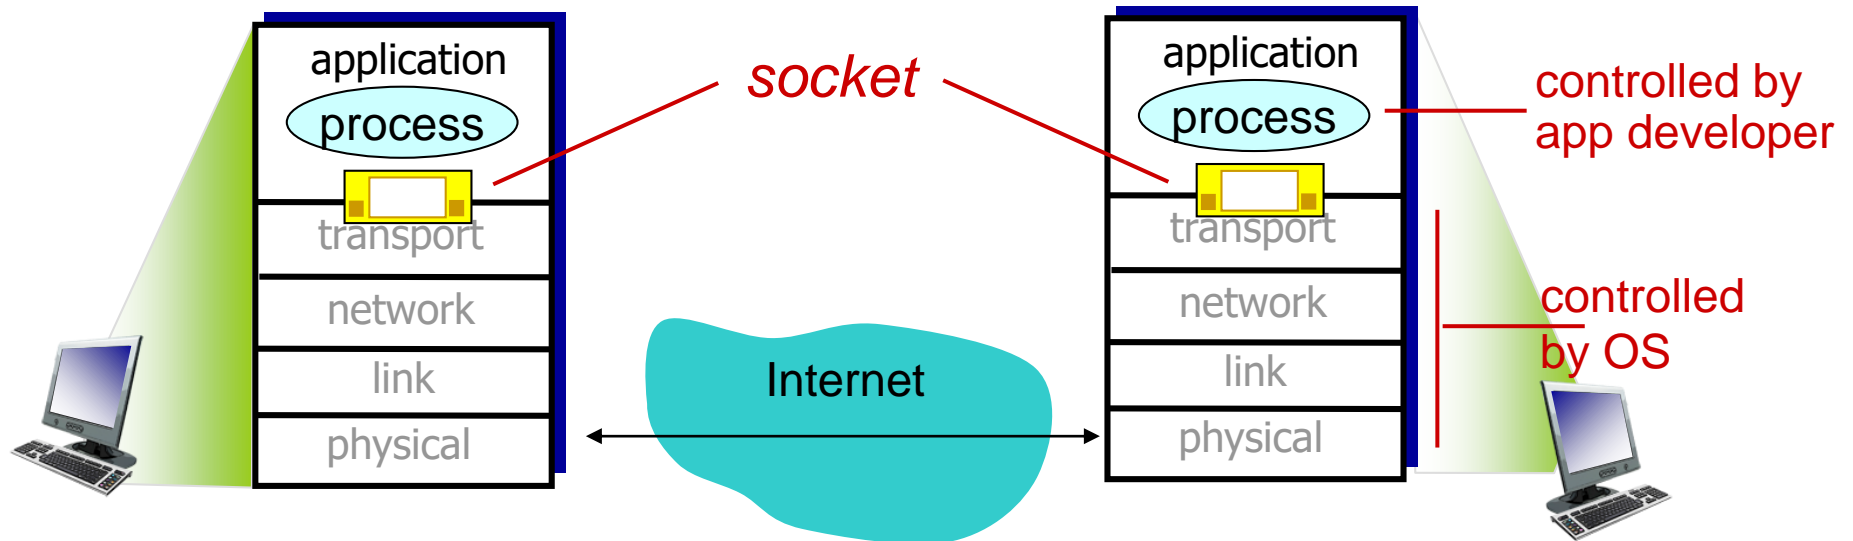
- Both sides need to have telephones installed
- Phone number is assigned to both sides
- Ringer is turned on to listen for calls
- Caller lifts phone and dials a number
- Phone rings and the receiver picks up the phone
- Conversation
- Both sides hang up once the call is over

In the Computer Universe

- `socket()`: endpoint for communication
- `bind()`: assign a unique (telephone) number
- `listen()`: wait for a call
- `connect()`: dial a number
- `accept()`: receive a call
- `send()`, `recv()`: conversation
- `close()`: hang up

Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



Sockets

- A socket is a file descriptor that lets an application read/write data from/to the network

```
int socket(int domain, int type, int protocol);
```

```
int fd;          /* socket descriptor */
if ((fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

- `socket` returns an integer (socket descriptor)
 - `fd < 0` indicates that an error occurred
 - socket descriptors are similar to file descriptors
- `AF_INET`: associates a socket with the Internet protocol family
- `SOCK_STREAM`: selects the TCP protocol
- `SOCK_DGRAM`: selects the UDP protocol

Internet Addressing Data Structure

```
#include <netinet/in.h>

/* Internet address structure */
struct in_addr {
    u_long s_addr;          /* 32-bit IPv4 address */
};                          /* network byte ordered */

/* Socket address, Internet style. */
struct sockaddr_in {
    u_char  sin_family;    /* Address Family */
    u_short sin_port;      /* UDP or TCP Port# */
                                /* network byte ordered */
    struct in_addr sin_addr; /* Internet Address */
    char     sin_zero[8];    /* unused */
};
```

`sin_family = AF_INET` selects Internet address family

Byte Ordering

- Big Endian →
- Sun Solaris, PowerPC, ...
- Little Endian →
- i386, alpha, ...
- Network byte order = Big Endian

c[0]	c[1]	c[2]	c[3]
128	2	194	95
95	194	2	128

Internet Addressing Data Structure

- Converts between host byte order and network byte order
 - ‘h’ = host byte order
 - ‘n’ = network byte order
 - ‘l’ = long (4 bytes), converts IP addresses
 - ‘s’ = short (2 bytes), converts port numbers

```
#include <netinet/in.h>
```

```
unsigned long int htonl(unsigned long int hostlong);  
unsigned short int htons(unsigned short int hostshort);  
unsigned long int ntohl(unsigned long int netlong);  
unsigned short int ntohs(unsigned short int netshort);
```

Create Socket

- A socket is a file descriptor that lets an application read/write data from/to the network

```
int socket(int domain, int type, int protocol);
```

```
int fd;          /* socket descriptor */
if ((fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

- **socket** returns an integer (socket descriptor)
 - $fd < 0$ indicates that an error occurred
 - socket descriptors are similar to file descriptors
- **AF_INET**: associates a socket with the Internet protocol family
- **SOCK_STREAM**: selects the TCP protocol
- **SOCK_DGRAM**: selects the UDP protocol

Bind Socket to Address

- A socket can be bound to a port

```
int bind(int sockfd, const struct sockaddr
*addr, socklen_t addrlen);
```

```
int fd;                /* socket descriptor */
struct sockaddr_in srv; /* used by bind() */

/* create the socket */

srv.sin_family = AF_INET; /* use the Internet addr family */
srv.sin_port = htons(80); /* bind socket 'fd' to port 80*/

/* bind: a client may connect to any of my addresses */
srv.sin_addr.s_addr = htonl(INADDR_ANY);

if(bind(fd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("bind"); exit(1);
}
```

Listen to the port

- listen() indicates that the server will accept a connection

```
int listen(int sockfd, int backlog);
```

```
int fd;                                /* socket descriptor */
struct sockaddr_in srv;                 /* used by bind() */

/* 1) create the socket */
/* 2) bind the socket to a port */

if(listen(fd, 5) < 0) {
    perror("listen");
    exit(1);
}
```

Accept Connections

- `accept` returns a new socket (`newfd`) with the same properties as the original socket (`fd`)

```
int accept(int sockfd, struct sockaddr *addr,  
socklen_t *addrlen);
```

```
int fd;                /* socket descriptor */  
struct sockaddr_in srv; /* used by bind() */  
struct sockaddr_in cli; /* used by accept() */  
int newfd;              /* returned by accept() */  
int cli_len = sizeof(cli); /* used by accept() */  
  
/* 1) create the socket */  
/* 2) bind the socket to a port */  
/* 3) listen on the socket */  
  
newfd = accept(fd, (struct sockaddr*) &cli, &cli_len);  
if(newfd < 0) {  
    perror("accept"); exit(1);  
}
```

Read Data

- read blocks waiting for data from the client but does not guarantee that sizeof(buf) is read

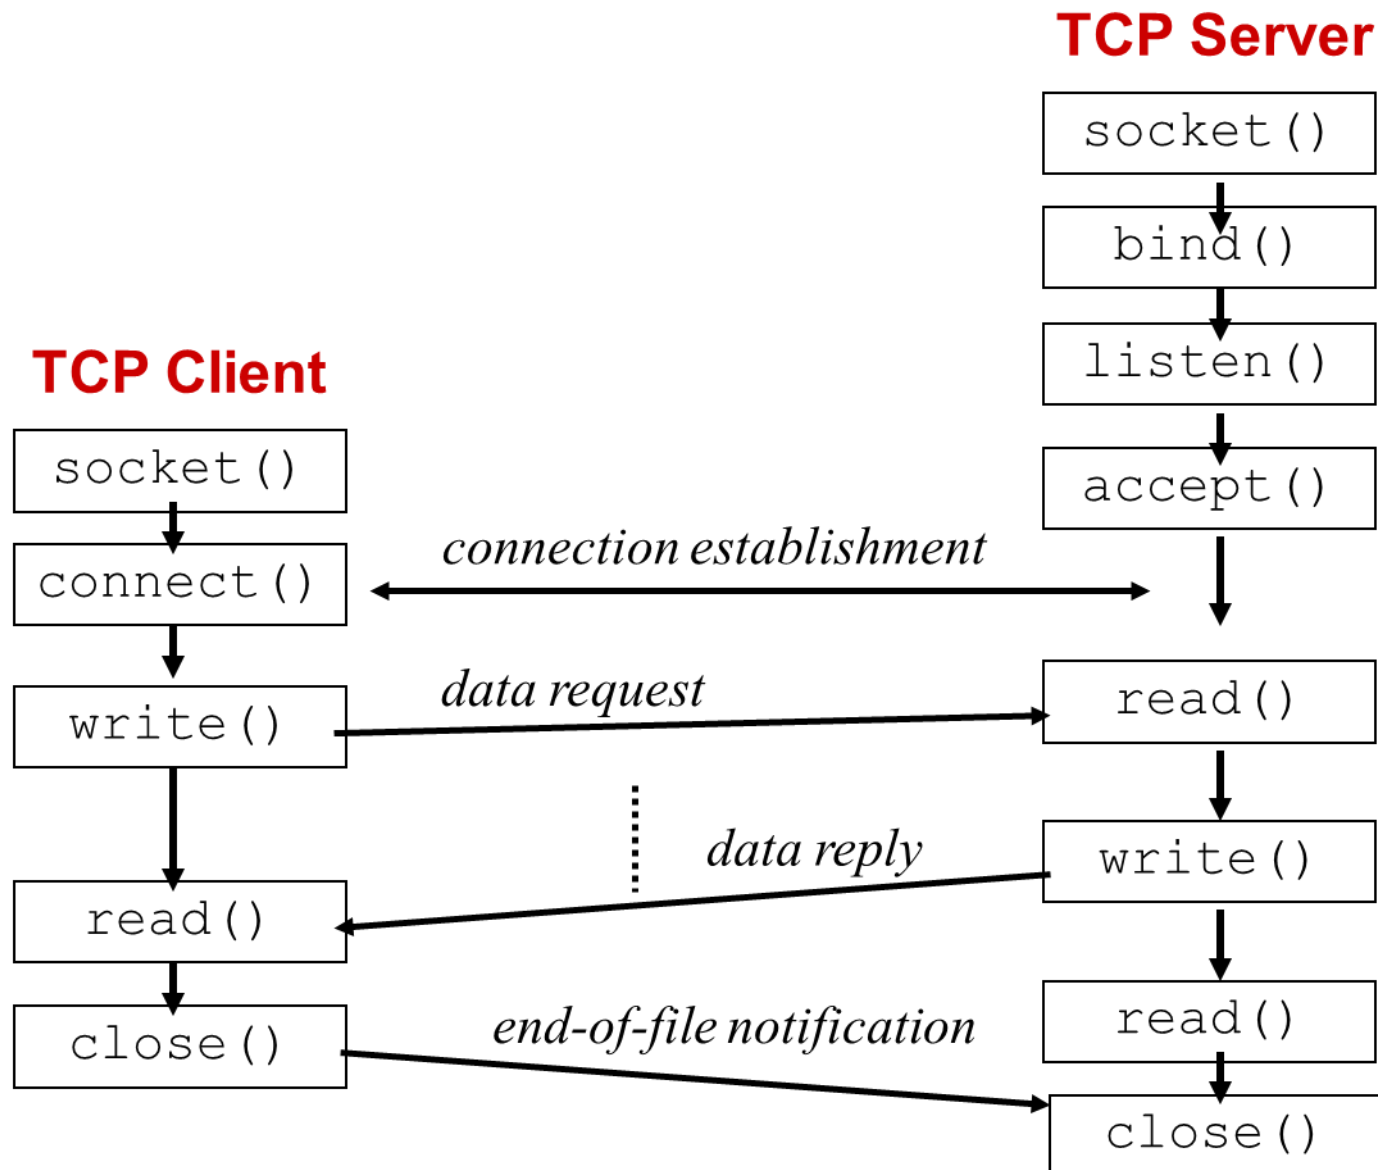
```
int read(int fd, void *buf, size_t count);
```

```
int fd;                /* socket descriptor */
char buf[512];          /* used by read() */
int nbytes;             /* used by read() */

/* 1) create the socket */
/* 2) bind the socket to a port */
/* 3) listen on the socket */
/* 4) accept the incoming connection */

if((nbytes = read(newfd, buf, sizeof(buf))) < 0) {
    perror("read"); exit(1);
}
```

Client



Dealing with IP Addresses

- IP Addresses are commonly written as strings (128.113.26.47), but programs deal with IP addresses as integers.

Converting strings to numerical address:

```
struct sockaddr_in srv;  
  
srv.sin_addr.s_addr = inet_addr("128.2.35.50") ;  
if(srv.sin_addr.s_addr == (in_addr_t) -1) {  
    fprintf(stderr, "inet_addr failed!\n");  
    exit(1);  
}
```

Converting a numerical address to a string:

```
struct sockaddr_in srv;  
char *t = inet_ntoa(srv.sin_addr) ;  
if(t == 0) {  
    fprintf(stderr, "inet_ntoa failed!\n");  
    exit(1);  
}
```

Dealing with IP Addresses

- gethostbyname provides interface to DNS
- Additional useful calls
 - gethostbyaddr : returns hostent given sockaddr_in
 - getservbyname: used to get service description (typically port number)

```
#include <netdb.h>

struct hostent *hp; /*ptr to host info for remote*/
struct sockaddr_in peeraddr;
char *name = "www.nus.edu.sg";

peeraddr.sin_family = AF_INET;
hp = gethostbyname(name)
peeraddr.sin_addr.s_addr = ((struct in_addr*) (hp->h_addr)) ->s_addr;
```

Connect to the Server

- `connect()` allows a client to connect to a server

```
int connect(int sockfd, const struct sockaddr
*addr, socklen_t addrlen);
```

```
int fd;                                /* socket descriptor */
struct sockaddr_in srv;                 /* used by connect() */

/* create the socket */

/* connect: use the Internet address family, and connect to
   port 80 on IP address "128.2.35.50" */
srv.sin_family = AF_INET;
srv.sin_port = htons(80);
srv.sin_addr.s_addr = inet_addr("128.2.35.50");

if(connect(fd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("connect"); exit(1);
}
```

Write the Data

- `write()` can be used with a socket

```
ssize_t write(int fd, const void *buf, size_t
count);
```

```
int fd;                /* socket descriptor */
struct sockaddr_in srv; /* used by connect() */
char buf[512];          /* used by write() */
int nbytes;             /* used by write() */

/* 1) create the socket */
/* 2) connect() to the server */

/* Example: A client could "write" a request to a server */
if((nbytes = write(fd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```

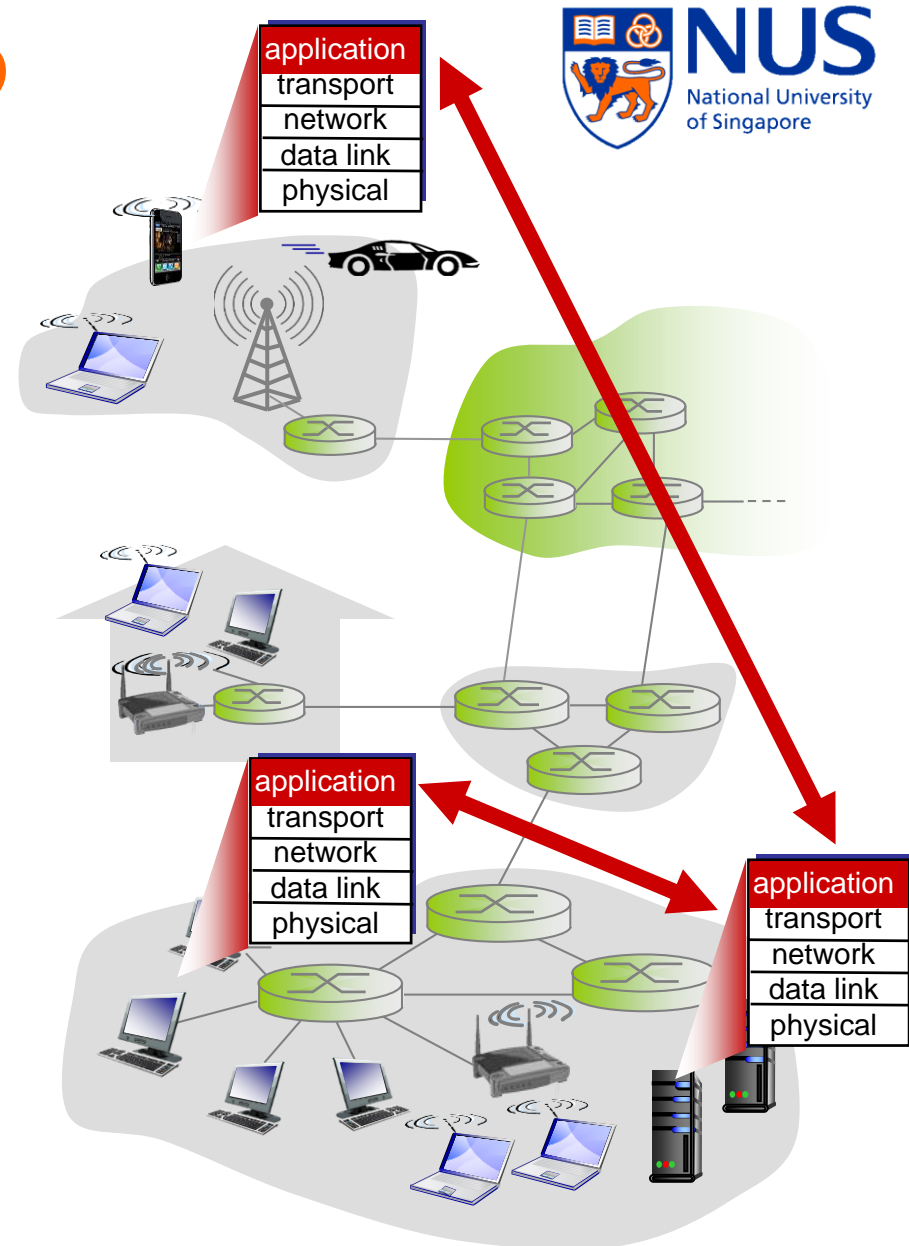
Creating a network app

write programs that:

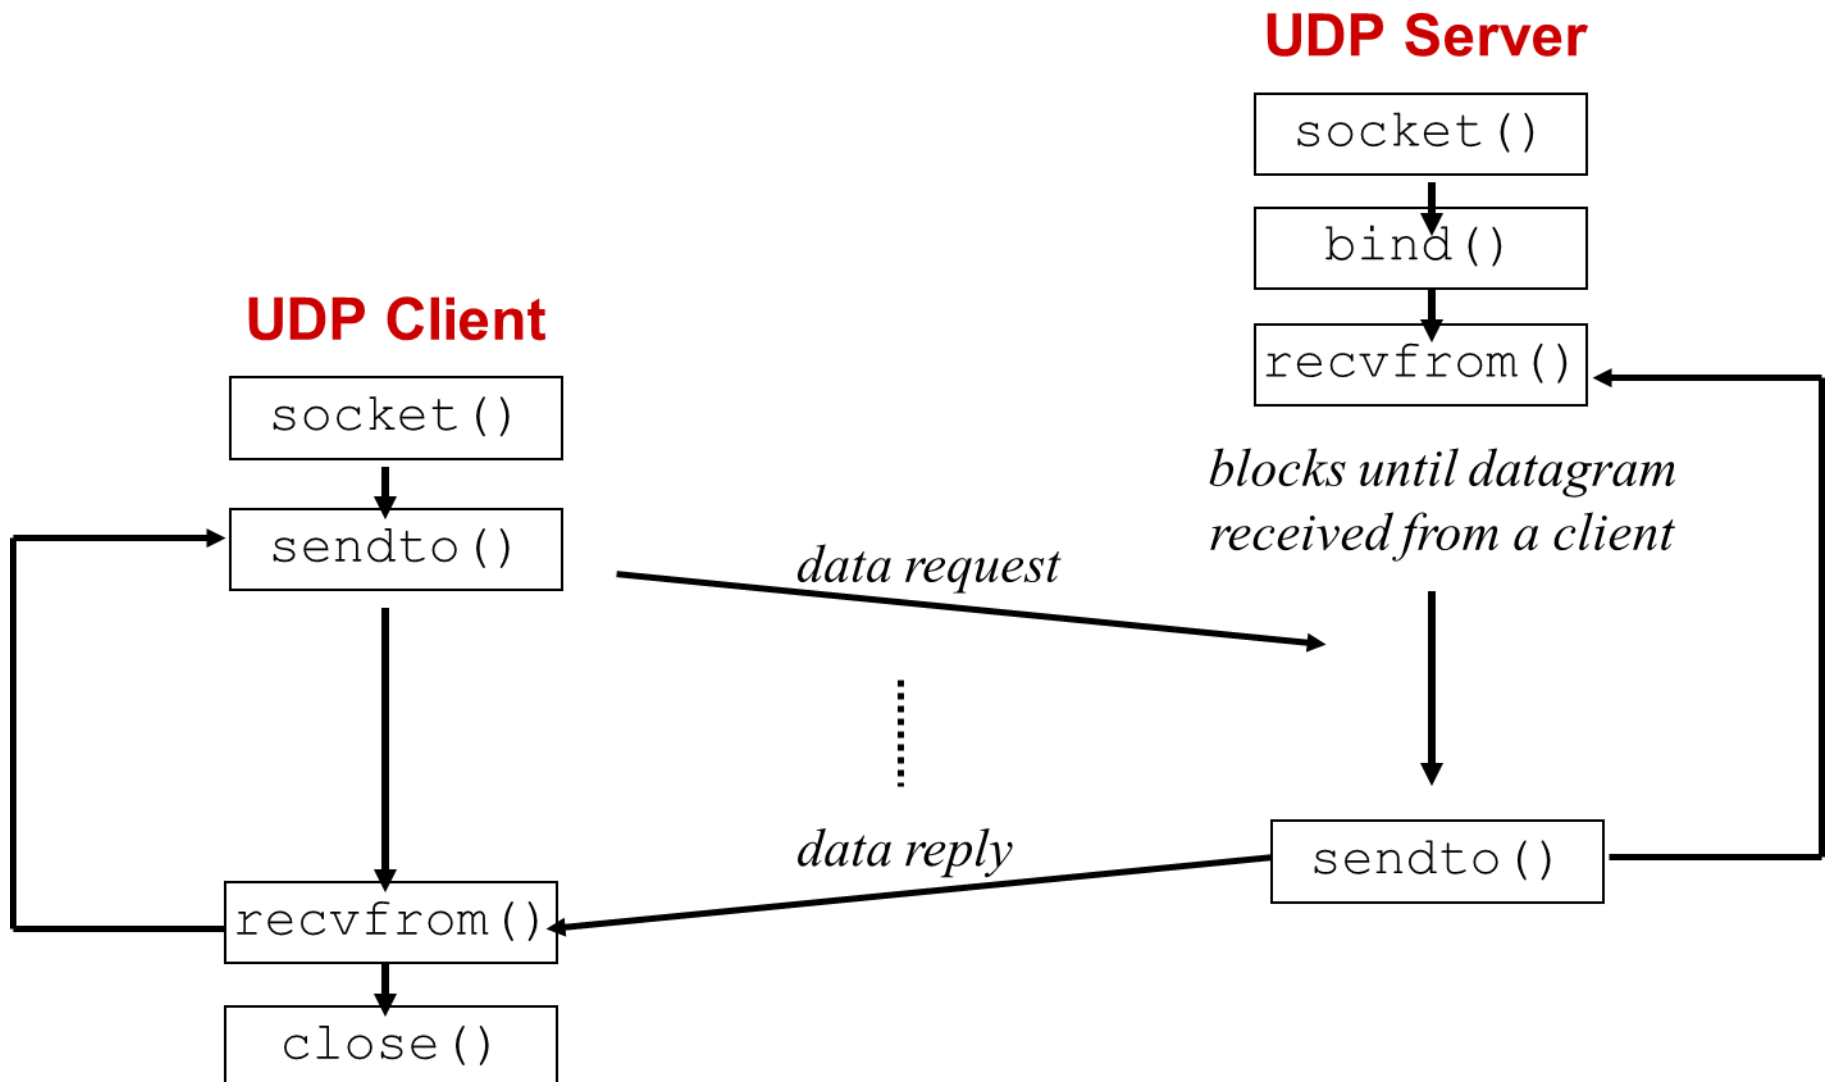
- run on (different) *end systems*
- communicate over network
- e.g., web server software communicates with browser software

no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation




UDP Socket Programming



UDP: Server

- What does a UDP server need to do so that a UDP client can connect to it?



Example: NTP daemon

Create a socket

- The UDP server must create a datagram socket
 - `SOCK_DGRAM`: selects the UDP protocol

```
int fd;          /* socket descriptor */  
  
if((fd = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {  
    perror("socket");  
    exit(1);  
}
```


Bind the socket

```
int fd;                /* socket descriptor */
struct sockaddr_in srv; /* used by bind() */

/* create the socket */

/* bind: use the Internet address family */
srv.sin_family = AF_INET;

/* bind: socket 'fd' to port 80*/
srv.sin_port = htons(80);

/* bind: a client may connect to any of my addresses */
srv.sin_addr.s_addr = htonl(INADDR_ANY);

if(bind(fd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("bind"); exit(1);
}
```

Read Data

- `read()` does not provide the client's address to the UDP server

```
ssize_t recvfrom(int sockfd, void *buf, size_t len,  
int flags, struct sockaddr *src_addr, socklen_t  
*addrlen);
```


```
int fd; /* socket descriptor */  
struct sockaddr_in srv; /* used by bind() */  
struct sockaddr_in cli; /* used by recvfrom() */  
char buf[512]; /* used by recvfrom() */  
int cli_len = sizeof(cli); /* used by recvfrom() */  
int nbytes; /* used by recvfrom() */  
  
/* 1) create the socket */  
/* 2) bind to the socket */  
  
nbytes = recvfrom(fd, buf, sizeof(buf), 0, (struct  
sockaddr*) &cli, &cli_len);  
if(nbytes < 0) {  
    perror("recvfrom"); exit(1);  
}
```

Read Data

- `nbytes = recvfrom(fd, buf, sizeof(buf), 0, (struct sockaddr*) &cli, &cli_len);`
- The actions performed by `recvfrom`
 - returns the number of bytes read (`nbytes`)
 - copies `nbytes` of data into `buf`
 - returns the address of the client (`cli`)
 - returns the length of `cli` (`cli_len`)

UDP: Client

- How does a UDP client communicate with a UDP server?



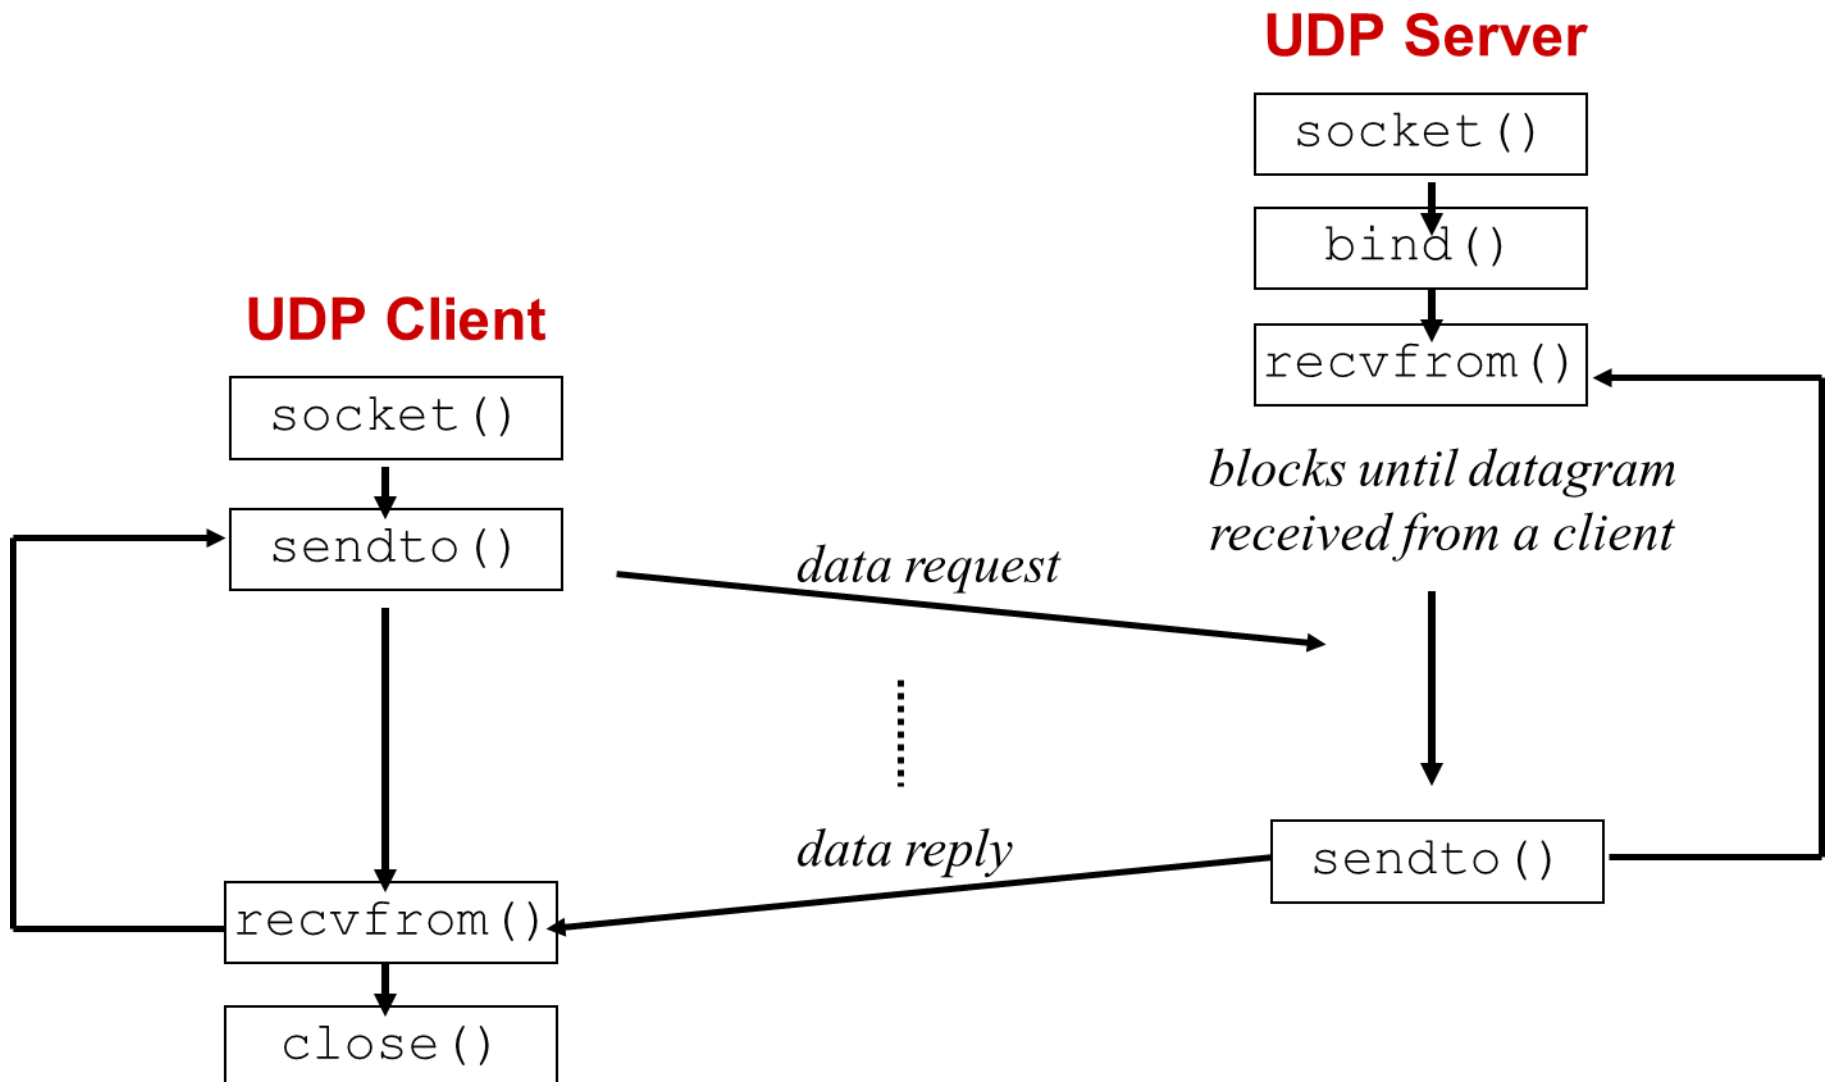
Send Data

- UDP client does not bind a port number: dynamically assigned when the first `sendto()` is called


```
ssize_t sendto(int sockfd, const void *buf, size_t  
len, int flags, const struct sockaddr *dest_addr,  
socklen_t addrlen);
```

```
int fd;                /* socket descriptor */  
struct sockaddr_in srv; /* used by sendto() */  
  
/* 1) create the socket */  
  
/* sendto: send data to IP Address "128.2.35.50" port 80 */  
srv.sin_family = AF_INET;  
srv.sin_port = htons(80);  
srv.sin_addr.s_addr = inet_addr("128.2.35.50");  
  
nbytes = sendto(fd, buf, sizeof(buf), 0 /* flags */,  
                (struct sockaddr*) &srv, sizeof(srv));  
if(nbytes < 0) {  
    perror("sendto"); exit(1);  
}
```

UDP Socket Programming



UDP Server: Multiple Connections



UDP Server: Multiple Connections

```
int s1;          /* socket descriptor 1 */
int s2;          /* socket descriptor 2 */

/* 1) create socket s1 */
/* 2) create socket s2 */
/* 3) bind s1 to port 2000 */
/* 4) bind s2 to port 3000 */

while(1) {
    recvfrom(s1, buf, sizeof(buf), ...);
    /* process buf */

    recvfrom(s2, buf, sizeof(buf), ...);
    /* process buf */
}
```

- What problems does this code have?

UDP Server: Multiple Connections

```
int select(int maxfds, fd_set *readfds, fd_set *writelfds,  
          fd_set *exceptfds, struct timeval *timeout);  
  
FD_CLR(int fd, fd_set *fds);    /* clear the bit for fd in fds */  
FD_ISSET(int fd, fd_set *fds); /* is the bit for fd in fds? */  
FD_SET(int fd, fd_set *fds);    /* turn on the bit for fd in fds */  
FD_ZERO(fd_set *fds);          /* clear all bits in fds */
```

- **maxfds**: number of descriptors to be tested
 - descriptors (0, 1, ... maxfds-1) will be tested
- **readfds**: a set of fds we want to check if data is available
 - returns a set of fds ready to read
 - if input argument is NULL, not interested in that condition
- **writelfds**: returns a set of fds ready to write
- **exceptfds**: returns a set of fds with exception conditions

UDP Server: Multiple Connections

```
int select(int maxfds, fd_set *readfds, fd_set *writefds,  
          fd_set *exceptfds, struct timeval *timeout);  
  
struct timeval {  
    long tv_sec;           /* seconds */  
    long tv_usec;         /* microseconds */  
}
```

- **timeout**
 - if NULL, wait forever and return only when one of the descriptors is ready for I/O
 - otherwise, wait up to a fixed amount of time specified by timeout
- if we don't want to wait at all, create a timeout structure with timer value equal to 0

UDP Server: Multiple Connections

- Use `select()` for synchronous I/O multiplexing


```
int s1, s2;                /* socket descriptors */
fd_set readfds;            /* used by select() */

/* create and bind s1 and s2 */
while(1) {
    FD_ZERO(&readfds);      /* initialize the fd set */
    FD_SET(s1, &readfds);  /* add s1 to the fd set */
    FD_SET(s2, &readfds);  /* add s2 to the fd set */

    if(select(s2+1, &readfds, 0, 0, 0) < 0) {
        perror("select");
        exit(1);
    }
    if(FD_ISSET(s1, &readfds)) {
        recvfrom(s1, buf, sizeof(buf), ...);
        /* process buf */
    }
    /* do the same for s2 */
}
```

Web Servers

- How can a web server manage multiple connections simultaneously?



Web Server: Multiple Connections

- Use `select()` for synchronous I/O multiplexing

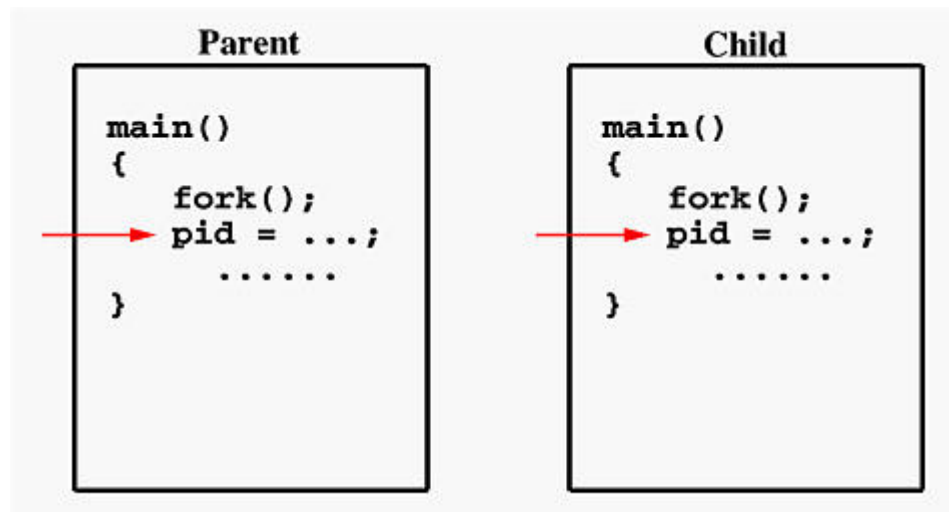
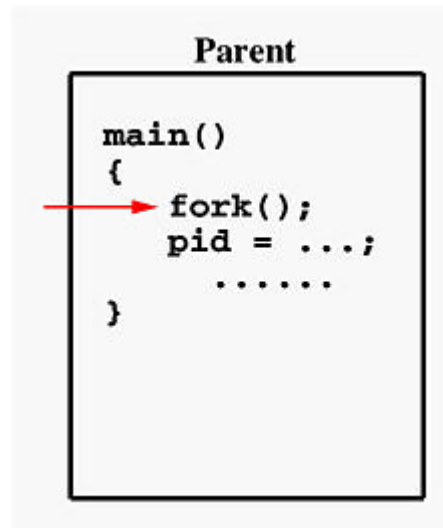
```
int fd, next=0;                /* original socket */
int newfd[10];                 /* new socket descriptors */
while(1) {
    fd_set readfds;
    FD_ZERO(&readfds); FD_SET(fd, &readfds);

    /* Now use FD_SET to initialize other newfd's
       that have already been returned by accept() */

    select(maxfd+1, &readfds, 0, 0, 0);
    if(FD_ISSET(fd, &readfds)) {
        newfd[next++] = accept(fd, ...);
    }
    /* do the following for each descriptor newfd[n] */
    if(FD_ISSET(newfd[n], &readfds)) {
        read(newfd[n], buf, sizeof(buf));
        /* process data */
    }
}
```

Concurrent Servers

- Use `fork()` for creating new (identical) processes



Concurrent Servers

- The fork() system call:

```
pid_t fork(void);
```

- If fork() returns a negative value, the creation of a child process was unsuccessful.
- fork() returns a zero to the newly created child process.
- fork() returns a positive value, the process ID of the child process, to the parent.
- The returned process ID is of type pid_t defined in sys/types.h.

Server Outline

```
pid_t pid;
int listenfd, connfd;

listenfd = socket(. . .);
bind(. . .);
listen(. . .);

while(1) {
    connfd = accept(listenfd, . . .);

    if ((pid = fork()) == 0) {
        close(listenfd);
        ...
        close(connfd);
        exit(0);
    }

    close(connfd);
}
```


Socket Programming References

- Man page
 - usage: man <function name>
- Textbook (Kurose and Ross)
 - Section 2.7
 - demo programs written in Python
- Unix Network Programming : Networking APIs: Sockets and XTI (Volume 1)
 - ultimate socket programming reference

TCP State Diagram

